

▼ Credit Card Churn Prediction Based on Spending Behaviors

0. Setup Google Drive Environment and Get Data

```
!pip install -U -q PyDrive

from pydrive.auth import GoogleAuth
from pydrive.drive import GoogleDrive
from google.colab import auth
from oauth2client.client import GoogleCredentials

auth.authenticate_user()
gauth = GoogleAuth()
gauth.credentials = GoogleCredentials.get_application_default()
drive = GoogleDrive(gauth)

file = drive.CreateFile({'id':'1Nwnmneol2ApYRvl8uBYtFsNNbSDVWiz6'}) # replace the id with id of file you want to access
file.GetContentFile('data.csv')

import pandas as pd
import numpy as np

df = pd.read_csv('data.csv')
df.head()
```

	CLIENTNUM	Attrition_Flag	Customer_Age	Gender	Dependent_count	Education_Level	Marital_Status
0	768805383	Existing Customer	45	M	3	High School	Married
1	818770008	Existing Customer	49	F	5	Graduate	Single
2	713982108	Existing Customer	51	M	3	Graduate	Married
3	769911858	Existing Customer	40	F	4	High School	Unknown
4	709106358	Existing Customer	40	M	3	Uneducated	Married

5 rows × 23 columns



1. Data Overview

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10127 entries, 0 to 10126
Data columns (total 23 columns):
#   Column                                     Non-Null
---  -----
0   CLIENTNUM                                10127 n
1   Attrition_Flag                           10127 n
2   Customer_Age                             10127 n
3   Gender                                   10127 n
4   Dependent_count                          10127 n
5   Education_Level                          10127 n
6   Marital_Status                           10127 n
7   Income_Category                          10127 n
8   Card_Category                             10127 n
9   Months_on_book                           10127 n
10  Total_Relationship_Count                  10127 n
11  Months_Inactive_12_mon                    10127 n
12  Contacts_Count_12_mon                     10127 n
13  Credit_Limit                              10127 n
14  Total_Revolving_Bal                       10127 n
15  Avg_Open_To_Buy                           10127 n
16  Total_Amt_Chng_Q4_Q1                      10127 n
17  Total_Trans_Amt                           10127 n
18  Total_Trans_Ct                             10127 n
19  Total_Ct_Chng_Q4_Q1                       10127 n
20  Avg_Utilization_Ratio                     10127 n
21  Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_1 10127 n
```

```
2023/3/14 20:06 My Supervised Learning.ipynb - Colaboratory
22 Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_2 10127 n
dtypes: float64(7), int64(10), object(6)
memory usage: 1.8+ MB

df.nunique()

CLIENTNUM 10127
Attrition_Flag 2
Customer_Age 45
Gender 2
Dependent_count 6
Education_Level 7
Marital_Status 4
Income_Category 6
Card_Category 4
Months_on_book 44
Total_Relationship_Count 6
Months_Inactive_12_mon 7
Contacts_Count_12_mon 7
Credit_Limit 6205
Total_Revolving_Bal 1974
Avg_Open_To_Buy 6813
Total_Amt_Chng_Q4_Q1 1158
Total_Trans_Amt 5033
Total_Trans_Ct 126
Total_Ct_Chng_Q4_Q1 830
Avg_Utilization_Ratio 964
Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_1 1704
Naive_Bayes_Classifier_Attrition_Flag_Card_Category_Contacts_Count_12_mon_Dependent_count_Education_Level_Months_Inactive_12_mon_2 640
dtype: int64
```

As shown above, this data set has 10127 rows and 22 columns from 4 categories:

**(1).Demographic Information**

- CLIENTNUM: Unique identifier for each customer.
- Customer\_Age: Age of customer.
- Gender: Gender of customer.
- Dependent\_count: Number of dependents that customer has.
- Education\_Level: Education level of customer.
- Marital\_Status: Marital status of customer.
- Income\_Category: Income category of customer.

**(2).Relationship with Card Provider**

- Card\_Category: Type of card held by customer.
- Months\_on\_book: How long customer has been on the books.
- Total\_Relationship\_Count: Total number of relationships customer has with the credit card provider.
- Months\_Inactive\_12\_mon: Number of months customer has been inactive in the last twelve months.
- Contacts\_Count\_12\_mon: Number of contacts customer has had in the last twelve months.
- Credit\_Limit: Credit limit of customer.

**(3).Spending Behavior**

- Total\_Revolving\_Bal: Total revolving balance of customer.
- Avg\_Open\_To\_Buy: Average open to buy ratio of customer.
- Total\_Amt\_Chng\_Q4\_Q1: Total amount changed from quarter 4 to quarter 1.
- Total\_Trans\_Amt: Total transaction amount.
- Total\_Trans\_Ct: Total transaction count.
- Total\_Ct\_Chng\_Q4\_Q1: Total count changed from quarter 4 to quarter 1.
- Avg\_Utilization\_Ratio: Average utilization ratio of customer.

**(4).Information for Prediction**

- Attrition\_Flag: Flag indicating whether or not the customer has churned out.
- Naive\_Bayes\_Classifier\_Attrition\_Flag\_Card\_Category\_Contacts\_Count\_12\_mon\_Dependent\_count\_Education\_Level\_Months\_Inactive\_12\_mon: Naive Bayes classifier for predicting whether or not someone will churn

2. Spending Behavior Preprocessing

Pick out spending behavior features from data

```
df_sp = df[['Total_Revolving_Bal', 'Avg_Open_To_Buy', 'Total_Amt_Chng_Q4_Q1', 'Total_Trans_Amt', 'Total_Trans_Ct', 'Total_Ct_Chng_Q4_Q1', 'Avg_Utilization_Ratio']]
df_sp.head()
```

	Total_Revolving_Bal	Avg_Open_To_Buy	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	Total_Trans_Ct	Tot
0	777	11914.0	1.335	1144	42	
1	864	7392.0	1.541	1291	33	
2	0	3418.0	2.594	1887	20	
3	2517	796.0	1.405	1171	20	
4	0	4716.0	2.175	816	28	

```
y = df['Attrition_Flag']
y.replace({'Existing Customer': 0, 'Attrited Customer': 1}, inplace=True)
y.head()

0    0
1    0
2    0
3    0
4    0
Name: Attrition_Flag, dtype: int64
```

We can see that features in spending behaviors are all described by numeric values. Let's take a look at them.

```
df_sp.isnull().sum()

Total_Revolving_Bal    0
Avg_Open_To_Buy        0
Total_Amt_Chng_Q4_Q1    0
Total_Trans_Amt         0
Total_Trans_Ct          0
Total_Ct_Chng_Q4_Q1     0
Avg_Utilization_Ratio    0
dtype: int64
```

```
df_sp.describe()
```

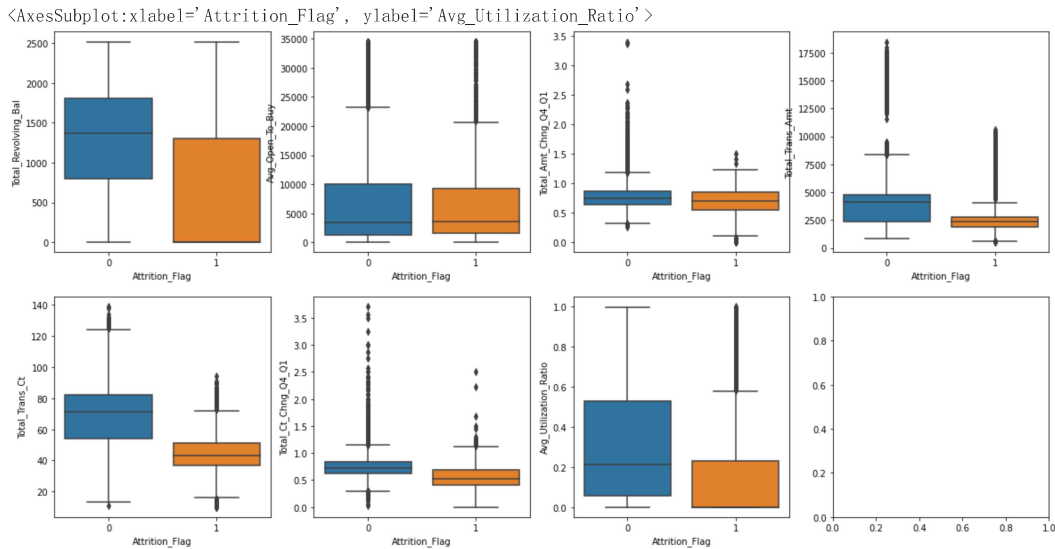
	Total_Revolving_Bal	Avg_Open_To_Buy	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	Total_Trans_Ct
count	10127.000000	10127.000000	10127.000000	10127.000000	10127.000000
mean	1162.814061	7469.139637	0.759941	4404.086304	64.858695
std	814.987335	9090.685324	0.219207	3397.129254	23.472570
min	0.000000	3.000000	0.000000	510.000000	10.000000
25%	359.000000	1324.500000	0.631000	2155.500000	45.000000
50%	1276.000000	3474.000000	0.736000	3899.000000	67.000000
75%	1784.000000	9859.000000	0.859000	4741.000000	81.000000
max	2517.000000	34516.000000	3.397000	18484.000000	139.000000

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df_plt = df_sp.join(y)
df_plt.head()
```

	Total_Revolving_Bal	Avg_Open_To_Buy	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	Total_Trans_Ct	Tot
0	777	11914.0	1.335	1144	42	
1	864	7392.0	1.541	1291	33	
2	0	3418.0	2.594	1887	20	
3	2517	796.0	1.405	1171	20	
4	0	4716.0	2.175	816	28	

```
# boxplot
_,axss = plt.subplots(2,4, figsize=[20,10])
sns.boxplot(x='Attrition_Flag', y='Total_Revolving_Bal', data=df_plt, ax=axss[0][0])
sns.boxplot(x='Attrition_Flag', y='Avg_Open_To_Buy', data=df_plt, ax=axss[0][1])
sns.boxplot(x='Attrition_Flag', y='Total_Amt_Chng_Q4_Q1', data=df_plt, ax=axss[0][2])
sns.boxplot(x='Attrition_Flag', y='Total_Trans_Amt', data=df_plt, ax=axss[0][3])
sns.boxplot(x='Attrition_Flag', y='Total_Trans_Ct', data=df_plt, ax=axss[1][0])
sns.boxplot(x='Attrition_Flag', y='Total_Ct_Chng_Q4_Q1', data=df_plt, ax=axss[1][1])
sns.boxplot(x='Attrition_Flag', y='Avg_Utilization_Ratio', data=df_plt, ax=axss[1][2])
```



We can see that there are a lot of outliers in the data. Let's drop all the outliers and fill them with the median value of corresponding features.

```
def iqr_outlier_rm(dt_input):
    lq,uq=np.percentile(dt_input,[25,75])
    lower_l=lq - 1.5*(uq-lq)
    upper_l=uq + 1.5*(uq-lq)
    return dt_input[(dt_input >= lower_l) & (dt_input <= upper_l)]

df_sp_ws = iqr_outlier_rm(df_sp)
#df_plt_ws.dropna(axis = 0, how = 'any', inplace = True)
df_sp_ws.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10127 entries, 0 to 10126
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Total_Revolving_Bal    10127 non-null  int64
1   Avg_Open_To_Buy        5662 non-null   float64
2   Total_Amt_Chng_Q4_Q1   10127 non-null  float64
3   Total_Trans_Amt        6470 non-null   float64
4   Total_Trans_Ct         10127 non-null  int64
5   Total_Ct_Chng_Q4_Q1    10127 non-null  float64
6   Avg_Utilization_Ratio  10127 non-null  float64
dtypes: float64(5), int64(2)
memory usage: 553.9 KB

df_sp_ws['Avg_Open_To_Buy'].fillna(df_sp_ws['Avg_Open_To_Buy'].median(), inplace = True)
df_sp_ws['Total_Trans_Amt'].fillna(df_sp_ws['Avg_Open_To_Buy'].median(), inplace = True)
df_sp_ws.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10127 entries, 0 to 10126
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Total_Revolving_Bal    10127 non-null  int64
1   Avg_Open_To_Buy        10127 non-null  float64
2   Total_Amt_Chng_Q4_Q1   10127 non-null  float64
```

```
3  Total_Trans_Amt      10127 non-null  float64
4  Total_Trans_Ct       10127 non-null  int64
5  Total_Ct_Chng_Q4_Q1   10127 non-null  float64
6  Avg_Utilization_Ratio 10127 non-null  float64
dtypes: float64(5), int64(2)
memory usage: 553.9 KB
```

```
df_sp_ws.describe()
```

	Total_Revolving_Bal	Avg_Open_To_Buy	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	Total_Trans_Ct
count	10127.000000	10127.000000	10127.000000	10127.000000	10127.000000
mean	1162.814061	1585.645463	0.759941	2237.537286	64.858695
std	814.987335	840.005231	0.219207	1050.033500	23.472570
min	0.000000	3.000000	0.000000	510.000000	10.000000
25%	359.000000	1324.500000	0.631000	1438.300000	45.000000
50%	1276.000000	1438.300000	0.736000	1658.000000	67.000000
75%	1784.000000	1521.500000	0.859000	3050.000000	81.000000
max	2517.000000	4391.000000	3.397000	4391.000000	139.000000

3. Model Training and Evaluation

(1) Preparation

Split data to train set and test set

```
from sklearn import model_selection

x_train, x_test, y_train, y_test = model_selection.train_test_split(df_sp_ws, y, test_size=0.25, stratify = y, random_state=1)
#I want to generate the same report everytime I run this colab
x_train_temp = x_train.copy()
x_train.head()
```

	Total_Revolving_Bal	Avg_Open_To_Buy	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	Total_Trans_Ct
9440	1583	1438.3	0.824	1438.3	112
959	1434	4127.0	1.423	1820.0	42
7737	2505	111.0	0.828	2576.0	42
7175	0	1438.3	0.809	2600.0	44
5844	0	1438.3	0.622	4333.0	84

Standarlize data sets

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
scaler.fit(x_train)
x_train = pd.DataFrame(scaler.transform(x_train))
x_test = pd.DataFrame(scaler.transform(x_test))
x_train.columns = df_sp.columns
x_test.columns = df_sp.columns
x_train.head()
```

	Total_Revolving_Bal	Avg_Open_To_Buy	Total_Amt_Chng_Q4_Q1	Total_Trans_Amt	Total_Trans_Ct	Tot
0	0.518958	-0.167468	0.292157	-0.761188	2.014456	
1	0.336068	3.037316	3.018093	-0.397718	-0.974368	
2	1.650670	-1.749538	0.310360	0.322175	-0.974368	
3	-1.424100	-0.167468	0.223895	0.345029	-0.888973	
4	-1.424100	-0.167468	-0.627106	1.995260	0.818926	

(2)Random Forest model

Train model

```

from sklearn.ensemble import RandomForestClassifier

classifier_RF = RandomForestClassifier()
classifier_RF.fit(x_train, y_train)
classifier_RF.predict(x_test)
classifier_RF.score(x_test, y_test)

0.938783570300158

```

### Use Grid Search to Find Optimal Hyperparameters Using Cross Validation(CV)

```

from sklearn.model_selection import GridSearchCV

# helper function for printing out grid search results
def print_grid_search_metrics(gs):
    print ("Best score: " + str(gs.best_score_))
    print ("Best parameters set:")
    best_parameters = gs.best_params_
    for param_name in sorted(best_parameters.keys()):
        print(param_name + ':' + str(best_parameters[param_name]))

parameters = {
    'n_estimators': [60, 80, 100],
    'max_depth': [1, 5, 10]
}
Grid_RF = GridSearchCV(RandomForestClassifier(), parameters, cv=5)
Grid_RF.fit(x_train, y_train)
print_grid_search_metrics(Grid_RF)

Best score: 0.9324555628703095
Best parameters set:
max_depth:10
n_estimators:80

# best random forest model
best_RF_model = Grid_RF.best_estimator_
best_RF_model.score(x_test, y_test)

0.9293048973143759

```

### ROC & AUC Evaluation of Random Forest Model


```

from sklearn.metrics import roc_curve
from sklearn import metrics

y_pred_rf = best_RF_model.predict_proba(x_test)[:, 1]
fpr_rf, tpr_rf, _ = roc_curve(y_test, y_pred_rf)

plt.figure(1)
plt.plot([0, 1], [0, 1], 'k--') # the diagonal
plt.plot(fpr_rf, tpr_rf, label='RF')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.title('ROC curve - RF model')
plt.legend(loc='best')
plt.show()

```



```

metrics.auc(fpr_rf, tpr_rf)

0.9672276340511635

```

Feature importance of Random Forest Model


```
importances = best_RF_model.feature_importances_  
indices = np.argsort(importances)[::-1]  
  
# Print the feature ranking  
print("Feature importance ranking by Random Forest Model:")  
for ind in range(len(indices)):  
    print ("{0} : {1}".format(x_train_temp.columns[indices[ind]],round(importances[indices[ind]], 4)))  
  
Feature importance ranking by Random Forest Model:  
Total_Trans_Ct : 0.2391  
Total_Trans_Amt : 0.1989  
Total_Ct_Chng_Q4_Q1 : 0.1659  
Total_Revolving_Bal : 0.1568  
Total_Amt_Chng_Q4_Q1 : 0.1015  
Avg_Utilization_Ratio : 0.0986  
Avg_Open_To_Buy : 0.0392
```

permutation Importance of Random Forest Model

```
from sklearn.inspection import permutation_importance  
  
PI_RF = permutation_importance(best_RF_model, x_test, y_test, n_repeats=5, random_state=1)  
PI_res = pd.DataFrame(data=np.transpose([PI_RF['importances_mean'],PI_RF['importances_std']]),  
                      index = x_train.columns,columns=['PI_mean','PI_std'])  
PI_res = PI_res.sort_values(by='PI_mean',ascending=False)  
PI_res
```

	PI_mean	PI_std
Total_Trans_Ct	0.074092	0.003458
Total_Trans_Amt	0.050474	0.002907
Total_Revolving_Bal	0.049131	0.006401
Total_Amt_Chng_Q4_Q1	0.015008	0.001172
Total_Ct_Chng_Q4_Q1	0.013586	0.003743
Avg_Utilization_Ratio	0.010664	0.002902
Avg_Open_To_Buy	0.002212	0.001134

```
plt_PI = sns.barplot(x="PI_mean", y=PI_res.index, data=PI_res)  
plt_PI.figure.set_size_inches(12, 9)
```



(3)K Nearest Neighbor model

**Train model**

```
from sklearn.neighbors import KNeighborsClassifier

classifier_KNN = KNeighborsClassifier()
classifier_KNN.fit(x_train, y_train)
classifier_KNN.score(x_test, y_test)

0.9210110584518167
```

**Use Grid Search to Find Optimal Hyperparameters Using Cross Validation(CV)**

```
parameters = {
    'n_neighbors': [1, 3, 5, 7, 9]
}
Grid_KNN = GridSearchCV(KNeighborsClassifier(), parameters, cv=5)
Grid_KNN.fit(x_train, y_train)
print_grid_search_metrics(Grid_KNN)

Best score: 0.9208689927583936
Best parameters set:
n_neighbors:7


best_KNN_model = Grid_KNN.best_estimator_
best_KNN_model.score(x_test, y_test)

0.9261453396524486
```

**ROC & AUC Evaluation of K Nearest Neighbor Model**

```
y_pred_rf = best_KNN_model.predict_proba(x_test)[:, 1]
fpr_rf, tpr_rf, _ = roc_curve(y_test, y_pred_rf)

plt.figure(1)
plt.plot([0, 1], [0, 1], 'k--') # the diagonal
plt.plot(fpr_rf, tpr_rf, label='RF')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.title('ROC curve - RF model')
plt.legend(loc='best')
plt.show()
```



```
metrics.auc(fpr_rf, tpr_rf)

0.9502078335019513
```

**(4) Logistic Regression Model****Train model**

```
from sklearn.linear_model import LogisticRegression

classifier_LR = LogisticRegression()
classifier_LR.fit(x_train, y_train)
classifier_LR.score(x_test, y_test)

0.8783570300157978
```



## Use Grid Search to Find Optimal Hyperparameters Using Cross Validation(CV)

```

parameters = {
    'penalty':('l2','l1'),
    'C':(0.01, 0.05, 0.1, 0.2, 1)
}
Grid_LR = GridSearchCV(LogisticRegression(solver='liblinear'),parameters, cv=5)
Grid_LR.fit(x_train, y_train)
print_grid_search_metrics(Grid_LR)

Best score: 0.8812376563528638
Best parameters set:
C:0.1
penalty:l1

best_LR_model = Grid_LR.best_estimator_
best_LR_model.score(x_test, y_test)

0.8791469194312796

```


## ROC &amp; AUC Evaluation of Logistic Regression Model

```

y_pred_rf = best_LR_model.predict_proba(x_test)[:, 1]
fpr_rf, tpr_rf, _ = roc_curve(y_test, y_pred_rf)

plt.figure(1)
plt.plot([0, 1], [0, 1], 'k--') # the diagonal
plt.plot(fpr_rf, tpr_rf, label='RF')
plt.xlabel('False positive rate')
plt.ylabel('True positive rate')
plt.title('ROC curve - RF model')
plt.legend(loc='best')
plt.show()

```



```

metrics.auc(fpr_rf, tpr_rf)

0.8555502240208123

```